

# Embedded Systems Design with Platform FPGAs

## Chapter -5 Spatial Design

**Dr. Bassam Jamil**  
**Adopted and updated from**  
**Ron Sass and Andrew G. Schmidt**



## Chapter 5 — Spatial Design



## Outline of This Chapter

### 5.1 Principles of Parallelism

- 5.1.1 Granularity

- 5.1.2 Degree of Parallelism

- 5.1.3 Spatial Organizations

### 5.2 Identifying Parallelism

- 5.2.1 Ordering

- 5.2.2 Dependence

- 5.2.3 Uniform Dependence Vectors

### 5.3 Spatial Parallelism with Platform FPGAs

- 5.3.1 Parallelism within FPGA Hardware Cores

- 5.3.2 Parallelism within FPGA Designs

# Spatial Design

- Single vs. Multiple threads:
  - Designs with **a single thread** of control: each instruction or operation is executed fully **prior to** execution of the next.
  - **Spatial Design**:
    - allows for **multiple** threads of control to be executed **simultaneously**.
    - In computer science terms, this is more generally known as **concurrency**.
- Example: Consider performing a search over a large array of numbers.
  - linear search:  $O(n)$
  - binary search:
    - Single processor:  $O(\log n)$ .
    - fully parallel search implementation:  $O(1)$
- It may not always be feasible to implement such a large number of parallel operations,
  - it is worth considering how to design systems that are able to take advantage of the exceptional parallelism available within FPGAs.

## Outline of This Chapter

### 5.1 Principles of Parallelism

- 5.1.1 Granularity

- 5.1.2 Degree of Parallelism

- 5.1.3 Spatial Organizations

### 5.2 Identifying Parallelism

- 5.2.1 Ordering

- 5.2.2 Dependence

- 5.2.3 Uniform Dependence Vectors

### 5.3 Spatial Parallelism with Platform FPGAs

- 5.3.1 Parallelism within FPGA Hardware Cores



- 5.3.2 Parallelism within FPGA Designs

## 5.1 Principles of Parallelism

- Computation can be described as:
  - **Sequential**: The same components are reused over time and the results are the accumulation of a **sequence** of operations.
  - **Parallel**: Multiple instances of components are used and some operations may happen **concurrently**.
- Example (See Next Slide):
  - Consider a task T: that is composed of a set of eight subtasks
  - $T = \{t_0, t_1, \dots, t_7\}$ .
  - T is executed sequentially and with four units.
  - **Sequentially** on a single component
    - task T takes  $8\Delta$  to complete.
  - **Four units** only takes  $3\Delta$  to complete task T.
  - Speedup=2.1, but improved design needs  $4\times$  resources.

## Example

- Speed up=2.1
- Second design requires 4× the resources.



**“Why not move  $t_7$  to unit 3 in time slot  $[1\Delta, 2\Delta]$ ?** Then the application will complete 4× faster.

**Answer:** If the programmer who developed the application assumed that it would be executed sequentially then any parallel execution of T’s subtasks must respect this ordering.

**So fundamental issue of this chapter: discover the parallelism in the sequential software reference design.**

### 5.1.1. Granularity

- T(which is software reference design)is set of subtasks  $\{t_i\}$  .
- Two ways to define the sub-tasks.
  - Fine-grain** parallelism: subtasks  $t_i$  are relatively small, maybe the equivalent of a **few instructions**.
  - Coarse-grain** parallelism: relatively large subtasks, such as an **entire subroutine**.
- The following table summarizes the types and features of granularity based on how large the sub-tasks are.



## 5.1.1. Granularity

Level of Parallelism	$T = \{t_1, t_2, \dots\}$	Sub-tasks: $\{t_i\}$	Notes
Bit-level	<b>T</b> operates on scalar values	Bit fields: organize parallel <b>t</b> <sub>i</sub> 's to work on each field in parallel	Advantage: independent of the application Disadvantage: the number of bits in the scalar tends to limit the parallelism.
Instruction-level	Program	instructions	Advantage: increases parallelism of sequential code.
Iteration-level (data-level)	<b>T</b> operates on Vectors	Sequence of instruction on vectors	Example: inner product $c = \sum_i a_i \times b_i$
Thread-level	Parallelism explicitly in multiple threads	Each sub-task is a thread, which does not execute the same instructions	

## 5.1.2. Degree of Parallelism


- Degree of parallelism (DOP) can refer to
  - The number of concurrent operations at a single moment of time,
  - A time varying function over the entire application, or
  - The average of a time-varying function over some task T, which may not include the entire application.

### In the Figure:

- **Maximum DOP** is 6, which occurs between time 5–7.
- **Average DOP** is 3.
  - Average = Area/Time =  $21/7 = 3$





### What this tells us is that:

- if these are iteration-level subtasks, having more than six cores would be overkill.
- Likewise, if hardware resources were scarce, having four cores instead of six would achieve 87% of the performance with 66% of the resources.







## 5.1.3. Spatial Organizations

- Taxonomy is based on instruction stream (single/multiple) and data stream (single/multiple).
- **Single Instruction Stream, Single Data Stream (SISD)**
  - Single processing unit consumes a single sequence (stream) of data.
  - Used in many general purpose computers.
- **Single Instruction Stream, Multiple Data Stream (SIMD)**
  - Uses simple Processing Elements (PEs) that are all performing instructions from a **single instruction** stream in lock step.
- **Multiple Instruction Stream, Single Data Stream (MISD)**
  - Breaks up a single stream of instructions (or subinstructions) into independent operations to be executed in parallel.
  - These independent streams of instructions are then assigned to separate processors, and a single **stream of data is passed from processor to processor** via explicit communication channels.
- **Multiple Instruction Stream, Multiple Data Stream (MIMD)**
  - Multiple independent processors, each with its own memory and own instructions.
  - The processors might communicate via some shared memory hierarchy or via explicit communication channels.



# Spatial Organizations: Examples

- **SIMD** : used for special instruction sets added to commodity processors and Graphical Processor Units (GPU) used in video cards.
- Also, examples of **SIMD**
  - Vector processors, pipeline processors
  - Stream processors
  - Difference between vector and stream processor:
    - In vector processors, the programmer just had to follow a handful of guidelines and the [compiler](#) was primarily responsible for detecting and implementing vector instructions.
    - In stream processors, the [programmer](#) is usually tasked with explicitly identifying “kernels” in their code — each which has its own local memory and will ultimately be mapped to an independent processor.
- **MISD**:
  - “systolic architectures” because it resembled blood flowing along arteries and veins.
- **MIMD**
  - In MIMD, the hardware could execute separate instruction streams,
  - It is possible to replicate the same instruction stream on every processor. This programming model, called Single Program Multiple Data (SPMD) (or data parallel programming).
  - SPMD is productive way to write scalable parallel applications.



# Vector and Pipelined Parallelism


- Assume:

- Task  $T = \{t_0, t_1, \dots, t_{n-1}\}$
- each task buffers its outputs with a register, clocks this core at  $\Delta = \max_{t \in T} \{\text{delay}(t)\}$
- a sequence of inputs  $I$  that produces a sequence of outputs  $O = \langle o_0, o_1, \dots, o_{m-1} \rangle$
- then it takes  $n$  cycles to produce a result
- Throughput (outputs/cycle) =  $\frac{m}{m+n}$
- Assuming that  $n \ll m$ , then throughput approaches 1


- Embedded systems move and process large vectors, much like any computing system.

- By starting a new operation before the last one completes,

- we are effectively overlapping consecutive operations of  $T$ . The time-space diagram shown in the Figure illustrates this for four subtasks.
- Assuming that  $n \ll m$ , then the degree of parallelism is the number of threads



Pipeline parallelism.



## Control Parallelism (Irregular or ad hoc parallelism)

- Examines a single stream of instructions to find operations that can be performed across multiple components in parallel.
- Control parallelism can be utilized at different granularity.
  - We will examine statement and operation granularity.
- Consider the following program that computes the roots of the polynomial  $ax^2 + bx + c$  with the quadratic formula:

$$x_0, x_1 = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

- The sequential code would be:

(s1) descr = sqrt(b\*b-4\*a\*c)

(s2) denom = 2\*a


(s3) x0 = (-b+descr)/denom

(s4) x1 = (-b-descr)/denom

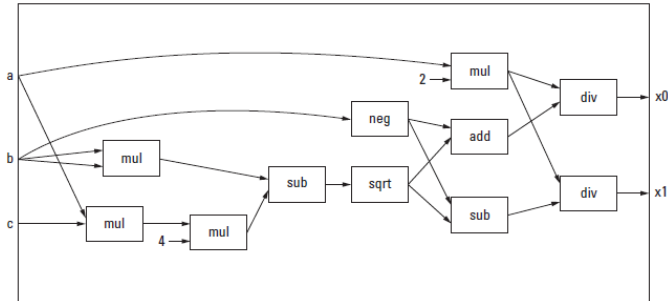
# Control Parallelism: Example

- Parallelism at **Statement Granularity**

- instructions s1 and s2 can operate in parallel,
- but s3 and s4 cannot begin until the first two finish.
- However, once s1 and s2 have finished, then the last two can be performed separately.
- the solution would have a constant degree of parallelism of two



- Parallelism at **Operation Granularity**



## Data (Regular) Parallelism

- Data parallelism focuses on organizing the parallel threads of control around a regular data structure, such as an array, and loops.
- If each iteration of a loop operating on an element of an array is independent, then it is **easy to divide up the array among the available resources**.
- In the event that the design is migrated to a **larger chip** with more resources, the array can be again divided by the available resources.



# Outline of This Chapter

## 5.1 Principles of Parallelism

- 5.1.1 Granularity

- 5.1.2 Degree of Parallelism

- 5.1.3 Spatial Organizations

## 5.2 Identifying Parallelism

- 5.2.1 Ordering

- 5.2.2 Dependence

- 5.2.3 Uniform Dependence Vectors

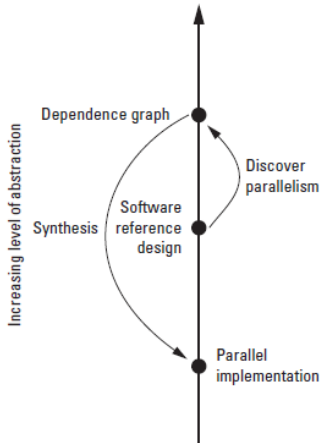
## 5.3 Spatial Parallelism with Platform FPGAs

- 5.3.1 Parallelism within FPGA Hardware Cores

- 5.3.2 Parallelism within FPGA Designs

## 5.2 Identifying Parallelism

- Finding parallelism is made more **difficult** because the **programming model** (which implements the software reference) is so closely associated with the **sequential compute model** of a typical processor.
- To implement parallelism from software model, we need to:
  - first move from a sequential specification to a higher level of abstraction
  - Second translate to a more concrete parallel implementation
- Nothing here is guaranteed to find the maximum parallelism in a software reference design.
  - To do that would require us to be able to infer a programmer's intent—a seemingly impossible to solve problem.



## 5.2.1 Ordering

- Assume a single task  $T$  composed of a sequence of  $n$  subtasks  $\langle t_0, t_1, \dots, t_{n-1} \rangle$ .
- We define the following terms.
- **Totally ordered** a composed sequence the subtasks in which  $t_i$  is started before  $t_{i+1}$  for all  $i$ .
- **Legal order** of the subtasks to be an ordering that produces the same results as total order  $T$  for all inputs.
- **Partial order** a binary relation that is reflexive, antisymmetric, and transitive; useful in this chapter to describe alternative legal orderings less strict than total orderings
- The goal of finding parallelism is to take a software reference design, and find a **legal partial order** from the expressed **total order**.

## Ordering: Example

- For example,

$a = (x+y) ; // (s1)$

$b = (x-y) ; // (s2)$

$r = a/b ; // (s3)$

- The following  $T$  and  $T'$  are both legal orderings.

–  $T = \langle s1, s2, s3 \rangle$

–  $T' = \langle s2, s1, s3 \rangle$

- Partial order:

– “ $s1$  has to complete before  $s3$ ” and

– “ $s2$  has to complete before  $s3$ ”

– No specific order is imposed for  $s1$  and  $s2$

## 5.2.2 Dependence

- Defining total order mathematically:

- Define a relation  $R$  on the set  $T$ . That is,  $R \subseteq T \times T$ .
- $(t_i, t_j) \in R$  if subtask  $t_i$  **comes before** subtask  $t_j$ 
  - then the total order would be  $(t_i, t_j) \in R_{\text{tot}}$  if  $i < j$  or  $(\forall 0 \leq i < n \mid (t_i, t_{i+1}))$

- What we really want is:

- a precedence relation that is not defined by the sequence of operations in the sequential software reference design.
- $(t_i, t_j) \in R$  if subtask  $t_i$  **must complete** before subtask  $t_j$
- “**must complete**” means:
  - task  $t_i$  produces some result that must be communicated to task  $t_j$  before  $t_j$  completes
  - These results are communicated through state, which is essentially the register file and primary memory (RAM) of a processor.
- So, precedence relation:  $(t_i, t_j) \in R$  if subtask  $t_i$  **writes a value that subtask  $t_j$  reads**

- Hence, The (true) dependence relation (i.e.,  $t_j$  depends on  $t_i$ ) is:

$$(t_i, t_j) \in D \quad \text{if} \quad \text{OUT}(t_i) \cap \text{IN}(t_j) \neq \emptyset \quad \text{where}$$

- $\text{OUT}(t_i)$  the values that  $t_i$  writes
- $\text{IN}(t_j)$  the values that  $t_j$  reads

## Dependency Example (1)

- The example experience:

- **(True) dependence** : s7 IN set includes **r**, which is in the s5 OUT set.
- **Output dependence**: s4 is output dependent on s6, as s5 has to complete before **a** can be overwritten.
- **Antidependence**: the value stored in **a** has to be used by s5 before “overwritten” by s6.

```
a = (x+y) ;    // (s4)
r = a ;        // (s5)
a = (x-y) ;    // (s6)
r = r/b ;      // (s7)
```

- **Can be resolved by renaming variables**

- We focus on true dependency.

### Recall Compute Architecture:

**Anti-dependency (write-after-read (WAR))** occurs when an instruction requires a value that is later updated.

**Output dependency( write-after-write (WAW))** occurs when the ordering of instructions will affect the final output value of a variable.

# Dependency Example (2): Generate the Dependency Set and Graph

- Consider example

- Generate IN/OUT table:

subtask $t$	$IN(t)$	$OUT(t)$
$s1 \ a = (x+y);$	$\{x, y\}$	$\{a\}$
$s2 \ b = (x-y);$	$\{x, y\}$	$\{b\}$
$s3 \ r = a/b;$	$\{a, b\}$	$\{r\}$

- Compute  $OUT(t_i) \cap IN(t_j)$  for all  $i$  and  $j$  (except when  $i = j$ , as it does not make sense that a subtask would be dependent on itself)




	s1 OUT	s2 OUT	s3 OUT
s1 IN		$OUT(s2) \cap IN(s1) = \{\}$	$OUT(s2) \cap IN(s1) = \{\}$
s2 IN	$OUT(s1) \cap IN(s2) = \{\}$		$OUT(t_3) \cap IN(t_2) = \{\}$
s3 IN	$OUT(s1) \cap IN(s3) = \{a\}$	$OUT(s2) \cap IN(s3) = \{b\}$	



- Dependence relation is all of the nonempty sets:  $D = \{(s1, s3), (s2, s3)\}$



Graphically, the dependency is shown in this figure



# Parallelism in if/else and while-loop



## • Parallelism in HDL constructs:

- **If/else**: use the hardware notion of a multiplexer with the condition determining which values to propagate.
- **“while loop”**: presents a greater chance for parallelism due to loop unrolling.
- Recognizing "induction variables" relationship to the loop control variables often lifts data dependence restrictions and permits further optimization

```

while ( C(i+0) or C(i+1) or C(i+2) ) {
  if ( C(i+0) )
    S(i+0)
  if ( C(i+1) )
    S(i+1)
  if ( C(i+2) )
    S(i+2)
  i = i+2
}
  
```

```

for ( i=0; i<n; i++ ) {
  a[j] = b[k];
  j++;
  k+=2;
}
  
```

```

_Kii2 = k;
_Kii1 = j;
for ( i = 0; i<n; i++ ) {
  a[_Kii1+i] = b[_Kii2+i*2];
}
  
```



## 5.2.3 Uniform Dependence Vectors

- **Dependence in loops** can be:
  - Within iteration
  - Between iterations.
- So, the **loop dependencies** are:
  - Loop-independent dependence
    - Dependence is within the iteration
    - But iterations are independent
  - Loop-carried dependence
    - Dependence between the iterations

# Uniform Dependence Vectors: loop-independent dependence

- Dependence vectors
  - treats the body of the loop(s) as a point in space.
  - the bounds of the loop(s) define a boundary in space.
- Singly nested loop has points on a line

```
for( i=0 ; i<10 ; i++ ) {  
    a[i] = c * x[i] ;  
}
```



Figure 5.13. Singly nested loop iteration space.

- Doubly nested loop:
  - there are three rows and five columns.
  - Inner  $i$  loop bounds are not constants, the iteration space is not rectangular.
  - Also notice that inside the body that  $s2$  depends on statement  $s1$ . This dependence is repeated in every iteration but it doesn't impact the execution of independent iterations.

```
for( j=1 ; j<=3 ; j++ ) {  
    for( i=1 ; i<=(j+5) ; i++ ) {  
        c = x[j][j] ;                // (s1)  
        a[i][j] = x[i][j]/c ;        // (s2)  
    }  
}
```





Figure 5.14. Doubly nested loop iteration space.

# Uniform Dependence Vectors: loop-carried dependence

## • Loop-Carried independence: single nested

- In the below example: statement s3 depends on itself — but in a different iteration.
- We highlight this relation between iterations with arrows
- iteration  $i+1$  depending on iteration  $i$ . We illustrate this as  $i \rightarrow (i+1)$  for all  $i$  within the bounds of space.


```
sum = 0 ;  
for( i=0 ; i<10 ; i++ ) {  
    sum = sum + x[i] ; // (s3)  
}
```



## • Loop-carried independence: doubly nested

- The below example results in an iteration space with the dependence vectors
- In the rows: there is the opportunity for pipeline parallelism.
- dependence vectors in this iteration space can be summarized with the distance vector  $(1, 0)$ .
- This is because iteration  $(j+1, i)$  depends on iteration  $(j, i)$ —taking the point-by-point difference leaves  $(1, 0)$ .

```
for i=1 to 3  
  for j=1 to 5  
    x[i] = x[i] + samp[i][j]
```



## 5.3. Spatial Parallelism with Platform FPGA

- Platform FPGAs
  - offer a wide variety of configurable logic resources that can be used in **parallel**.
  - offer the ability to reduce the number of integrated circuits needed in a design by migrating the logic onto the FPGA.
- This section looks
  - Looks more closely at **how to exploit the inherent parallelism** within the FPGA.
  - highlight some of the issues commonly associated with parallel implementations

## Processor vs. FPGA

- The differences between a microprocessor (single-core sequential processor) and a Platform FPGA with respect to parallel processing:
  1. For tough design constraints, a single controller may not suffice.
  2. An FPGA
    - provides the capability to **interface with a multitude** of **external devices**
    - Perform **intermediate and parallel** computations.

## 5.3.1 Parallelism Within FPGA Hardware Cores

- Expressing parallelism in HDL:

- If/else statement


- for-loop statement: instantiating multiple cores to operate in parallel

- How to parallelize a loop with loop-carried dependency?

```
sum = 0 ;  
for( i=0 ; i<10 ; i++ ) {  
    sum = sum + x[i] ; // (s3)  
}
```

- Answer: parallelize operations in hierarchy fashion.

- Speedup =  $10/4 = 2.5$



# Control of Sequential Operations

- **Handshaking:**
  - Go/Done
  - Consumer/Producer:
    - A producer is generating data and a consumer is operating on newly generated data .
    - The consumer must indicate when it is **ready** for data, and the producer indicates when new data are **valid**.
- **FSM:** Statemachine consists of (at least) the following states:
  - **idle:** the FSM waits for some condition (perhaps the go signal) before proceeding to the running state where the operation is executed.
  - **running:** the running state can be expanded to include more than just one state; however, the state machine is never in more than one state at a time.
  - **done:** Upon completion, the FSM may signal any number of dependent components via a done signal.
- **Pipelining :** by turning each state in the state machine into individual process blocks (or separate components).
  - Implementing pipelined systems may be easier than sequential finite state machines because a designer may chain together components to build complex operations. This chain can be pipelined with the addition of registers between each component.

## 5.3.2 Parallelism within FPGA Designs

- How to construct systems to achieve parallelism within FPGA designs
- **Add multiple instances** of the same hardware core.
  - This is explicit parallelism; there are physically more components running in parallel.
  - However, in order to take advantage of this new compute power, we must divide the data set among the compute cores.
- At a finer granularity we can replicate **multiple instances** of a component **within a single hardware** core



# Chapter-In-Review

- started with basic understanding of parallel and serial
- described a number concept related to parallelism (DOP, granularity, etc.)
- described parallelism in hardware and software
- introduced some formal mathematics to help identify parallelism in a software reference design



## Chapter 5 Terms

**sequential** the same components are re-used over time and the the results are the accumulation of a sequence of operations

**parallel** multiple instances of components are used and some operations may happen concurrently

**totally ordered** a composed sequence of  $n$  subtasks  $\langle t_0, t_1, \dots, t_{n-1} \rangle$  in which  $t_i$  is started before  $t_{i+1}$  for all  $i$

**legal order** of subtasks to be an ordering that produces the same results as total order  $T$  for all inputs.

**partial order** a binary relation that is reflexive, antisymmetric, and transitive; useful in this chapter to describe alternative legal orderings less strict than total orderings



## Chapter 5 Terms

**spatial** in hardware designs multiple threads of control are able to operate simultaneously

**fine-grain** parallelism means that the subtasks  $t_i$  are relatively small, maybe the equivalent of a few instructions

**coarse-grain** parallelism refers to relatively large subtasks, such as an entire subroutine

**degree of parallelism (DOP)** can refer to: the number of concurrent operations at a single moment of time; a time-varying function over the entire application; or the average of a time-varying function over some task  $T$  which may not include the entire application

